

DEEP LEARNING ASSIGNMENT

Image Classification Using Convolutional Neural Networks

Akash.E
2018103005
CSE - 'P' Batch
19/04/2021

PROBLEM STATEMENT :

Pneumonia is an inflammatory condition of the lung affecting primarily the small air sacs known as alveoli. We build a CNN model to predict if a given Chest X-Ray image is 'Normal' or affected with 'Pneumonia'. The disease may be classified by where it was acquired, such as community- or hospital-acquired or healthcare-associated pneumonia.

DATASET DETAILS :

The dataset is organized into 3 folders (train, test, val) and contains subfolders for each image category (Pneumonia/Normal).

There are 5,863 X-Ray images (JPEG) and 2 categories (Pneumonia/Normal). Chest X-ray images (anterior-posterior) were selected from retrospective cohorts of pediatric patients of one to five years old.

For the analysis of chest x-ray images, all chest radiographs were initially screened for quality control by removing all low quality or unreadable scans.

MODULE LIST :

1. Fetching the Data/Visualizing the Data.
2. Data Pre-processing.
3. Building the Model
4. Training and Validation.
5. Results and Visualization

1. FETCHING THE DATA /VISUALIZING THE DATA :

A function is defined to fetch the dataset and store it in three separate arrays for train, test and validation. The corresponding path is specified to load the dataset into the arrays.

cv2.imread function is used to read the dataset and convert it to grayscale.

```

> MI

#Importing necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import keras
from keras.models import Sequential
from keras.layers import Dense, Conv2D, MaxPool2D, Flatten, Dropout, BatchNormalization
from keras.preprocessing.image import ImageDataGenerator
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from keras.callbacks import ReduceLROnPlateau
import cv2
import os

> MI

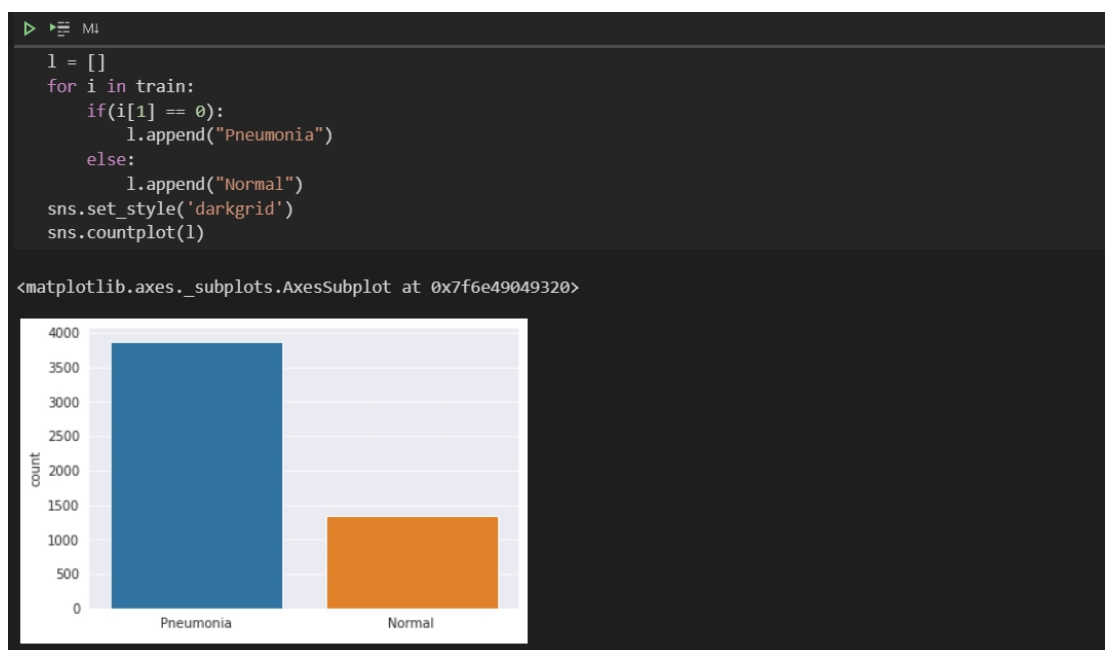
#Function to fetch the data
labels = ['PNEUMONIA', 'NORMAL']
img_size = 150
def get_training_data(data_dir):
    data = []
    for label in labels:
        path = os.path.join(data_dir, label)
        class_num = labels.index(label)
        for img in os.listdir(path):
            try:
                img_arr = cv2.imread(os.path.join(path, img), cv2.IMREAD_GRAYSCALE)
                resized_arr = cv2.resize(img_arr, (img_size, img_size)) # Reshaping images to preferred size
                data.append([resized_arr, class_num])
            except Exception as e:
                print(e)
    return np.array(data)

> MI

#Separate Arrays to store the data
train = get_training_data('../input/chest-xray-pneumonia/chest_xray/train')
test = get_training_data('../input/chest-xray-pneumonia/chest_xray/test')
val = get_training_data('../input/chest-xray-pneumonia/chest_xray/val')

```

Data population and Visualizing the data

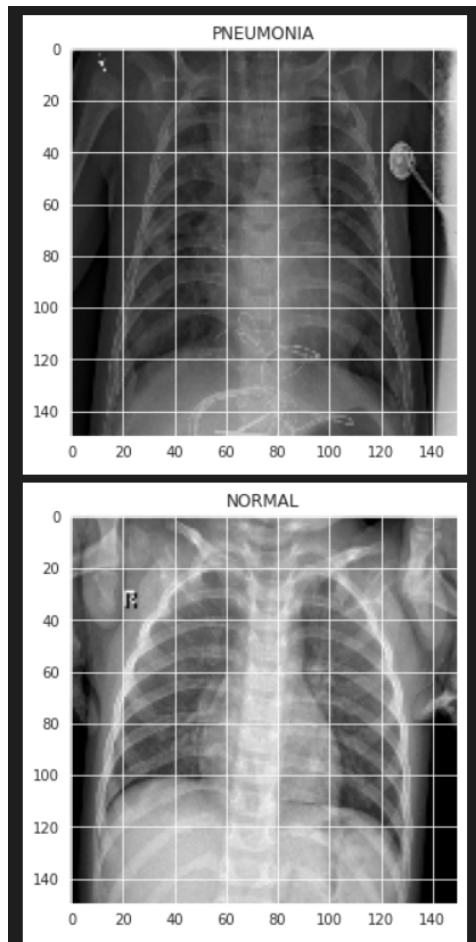


```

plt.figure(figsize = (5,5))
plt.imshow(train[0][0], cmap='gray')
plt.title(labels[train[0][1]])

plt.figure(figsize = (5,5))
plt.imshow(train[-1][0], cmap='gray')
plt.title(labels[train[-1][1]])

```



2. DATA PRE-PROCESSING :

Now the data is separated as x and y in the three arrays and preprocessing steps such as normalization and scaling occurs.

In order to avoid overfitting problem, we need to expand artificially our dataset. We can make your existing dataset even larger. The idea is to alter the training data with small transformations to reproduce the variations.

Approaches that alter the training data in ways that change the array representation while keeping the label the same are known as data augmentation techniques.

Some popular augmentations people use are grayscales, horizontal flips, vertical flips, random crops, color jitters, translations, rotations, and much more.

```

> ML
x_train = []
y_train = []

x_val = []
y_val = []

x_test = []
y_test = []

for feature, label in train:
    x_train.append(feature)
    y_train.append(label)

for feature, label in test:
    x_test.append(feature)
    y_test.append(label)

for feature, label in val:
    x_val.append(feature)
    y_val.append(label)

> ML
# Normalize the data
x_train = np.array(x_train) / 255
x_val = np.array(x_val) / 255
x_test = np.array(x_test) / 255

> ML
# resize data for deep learning
x_train = x_train.reshape(-1, img_size, img_size, 1)
y_train = np.array(y_train)

x_val = x_val.reshape(-1, img_size, img_size, 1)
y_val = np.array(y_val)

x_test = x_test.reshape(-1, img_size, img_size, 1)
y_test = np.array(y_test)

```

By applying just a couple of these transformations to our training data, we can easily double or triple the number of training examples and create a very robust model.

For data augmentation :

- Randomly rotate some training images by 30 degrees
- Randomly Zoom by 20% some training images
- Randomly shift images horizontally by 10% of the width
- Randomly shift images vertically by 10% of the height
- Randomly flip images horizontally. Once our model is ready, we fit the training dataset.

```

▶ MI
# With data augmentation to prevent overfitting and handling the imbalance in dataset

datagen = ImageDataGenerator(
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range = 30, # randomly rotate images in the range (degrees, 0 to 180)
    zoom_range = 0.2, # Randomly zoom image
    width_shift_range=0.1, # randomly shift images horizontally (fraction of total width)
    height_shift_range=0.1, # randomly shift images vertically (fraction of total height)
    horizontal_flip = True, # randomly flip images
    vertical_flip=False) # randomly flip images

datagen.fit(x_train)

```

3. BUILDING THE MODEL :

Now we build our CNN model with the following specifications :

```

▶ MI
model = Sequential()
model.add(Conv2D(32 , (3,3) , strides = 1 , padding = 'same' , activation = 'relu' , input_shape = (150,150,1)))
model.add(BatchNormalization())
model.add(MaxPool2D((2,2) , strides = 2 , padding = 'same'))
model.add(Conv2D(64 , (3,3) , strides = 1 , padding = 'same' , activation = 'relu'))
model.add(Dropout(0.1))
model.add(BatchNormalization())
model.add(MaxPool2D((2,2) , strides = 2 , padding = 'same'))
model.add(Conv2D(64 , (3,3) , strides = 1 , padding = 'same' , activation = 'relu'))
model.add(BatchNormalization())
model.add(MaxPool2D((2,2) , strides = 2 , padding = 'same'))
model.add(Conv2D(128 , (3,3) , strides = 1 , padding = 'same' , activation = 'relu'))
model.add(Dropout(0.2))
model.add(BatchNormalization())
model.add(MaxPool2D((2,2) , strides = 2 , padding = 'same'))
model.add(Conv2D(256 , (3,3) , strides = 1 , padding = 'same' , activation = 'relu'))
model.add(Dropout(0.2))
model.add(BatchNormalization())
model.add(MaxPool2D((2,2) , strides = 2 , padding = 'same'))
model.add(Flatten())
model.add(Dense(units = 128 , activation = 'relu'))
model.add(Dropout(0.2))
model.add(Dense(units = 1 , activation = 'sigmoid'))
model.compile(optimizer = "rmsprop" , loss = 'binary_crossentropy' , metrics = ['accuracy'])
model.summary()

```

Model summary :

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 150, 150, 32)	320
batch_normalization_1 (Batch Normalization)	(None, 150, 150, 32)	128
max_pooling2d_1 (MaxPooling2D)	(None, 75, 75, 32)	0
conv2d_2 (Conv2D)	(None, 75, 75, 64)	18496
dropout_1 (Dropout)	(None, 75, 75, 64)	0
batch_normalization_2 (Batch Normalization)	(None, 75, 75, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 38, 38, 64)	0
conv2d_3 (Conv2D)	(None, 38, 38, 64)	36928
batch_normalization_3 (Batch Normalization)	(None, 38, 38, 64)	256
max_pooling2d_3 (MaxPooling2D)	(None, 19, 19, 64)	0
conv2d_4 (Conv2D)	(None, 19, 19, 128)	73856
dropout_2 (Dropout)	(None, 19, 19, 128)	0
batch_normalization_4 (Batch Normalization)	(None, 19, 19, 128)	512
max_pooling2d_4 (MaxPooling2D)	(None, 10, 10, 128)	0
conv2d_5 (Conv2D)	(None, 10, 10, 256)	295168
dropout_3 (Dropout)	(None, 10, 10, 256)	0
batch_normalization_5 (Batch Normalization)	(None, 10, 10, 256)	1024
max_pooling2d_5 (MaxPooling2D)	(None, 5, 5, 256)	0
flatten_1 (Flatten)	(None, 6400)	0
dense_1 (Dense)	(None, 128)	819328
dropout_4 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 1)	129
Total params: 1,246,401		
Trainable params: 1,245,313		
Non-trainable params: 1,088		

4. TRAINING AND VALIDATION :

Now we train our data with the CNN we created. Reduce learning rate when a metric has stopped improving. Models often benefit from reducing the learning rate by a factor of 2-10 once learning stagnates.

This callback monitors a quantity and if no improvement is seen for a 'patience' number of epochs, the learning rate is reduced. For this we use **ReduceLROnPlateau** class.

Parameters :

Learning Rate	Default LR = 0.01 , reduced by a factor 0.3.
Epochs	5
Batch Size	32
Optimizer	Root Mean Square Prop - 'rmsprop'
Loss	Binary Cross Entropy

```
> *≡ ML
learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy', patience = 2, verbose=1, factor=0.3, min_lr=0.000001)

> *≡ ML
history = model.fit(datagen.flow(x_train,y_train, batch_size = 32), epochs = 5, validation_data = datagen.flow(x_val, y_val), callbacks = [learning_rate_reduction])

Epoch 1/5
163/163 [=====] - 281s 2s/step - loss: 0.3494 - accuracy: 0.8620 - val_loss: 9.9030 - val_accuracy: 0.5000
Epoch 2/5
163/163 [=====] - 250s 2s/step - loss: 0.2613 - accuracy: 0.9011 - val_loss: 14.1557 - val_accuracy: 0.5000
Epoch 3/5
163/163 [=====] - 250s 2s/step - loss: 0.2093 - accuracy: 0.9285 - val_loss: 25.8908 - val_accuracy: 0.5000

Epoch 00003: ReduceLROnPlateau reducing learning rate to 0.00030000000142492354.
Epoch 4/5
163/163 [=====] - 273s 2s/step - loss: 0.1464 - accuracy: 0.9519 - val_loss: 2.1655 - val_accuracy: 0.5625
Epoch 5/5
163/163 [=====] - 253s 2s/step - loss: 0.1305 - accuracy: 0.9536 - val_loss: 3.7105 - val_accuracy: 0.5000
```

Accuracy and Loss :

```
> *≡ ML

print("Loss of the model is - " , model.evaluate(x_test,y_test)[0])
print("Accuracy of the model is - " , model.evaluate(x_test,y_test)[1]*100 , "%")

624/624 [=====] - 9s 14ms/step
Loss of the model is - 0.3520233004521101
624/624 [=====] - 8s 13ms/step
Accuracy of the model is - 87.66025900840759 %
```


5. RESULTS AND VISUALIZATION :


Analysis after training the model :

```
epochs = [i for i in range(5)]
fig, ax = plt.subplots(1,2)
train_acc = history.history['accuracy']
train_loss = history.history['loss']
val_acc = history.history['val_accuracy']
val_loss = history.history['val_loss']
fig.set_size_inches(20,10)

ax[0].plot(epochs, train_acc, 'go-', label='Training Accuracy')
ax[0].plot(epochs, val_acc, 'ro-', label='Validation Accuracy')
ax[0].set_title('Training & Validation Accuracy')
ax[0].legend()
ax[0].set_xlabel("Epochs")
ax[0].set_ylabel("Accuracy")

ax[1].plot(epochs, train_loss, 'g-o', label='Training Loss')
ax[1].plot(epochs, val_loss, 'r-o', label='Validation Loss')
ax[1].set_title('Testing Accuracy & Loss')
ax[1].legend()
ax[1].set_xlabel("Epochs")
ax[1].set_ylabel("Training & Validation Loss")
plt.show()
```

Graph Plot :



Classification Report and Confusion Matrix :

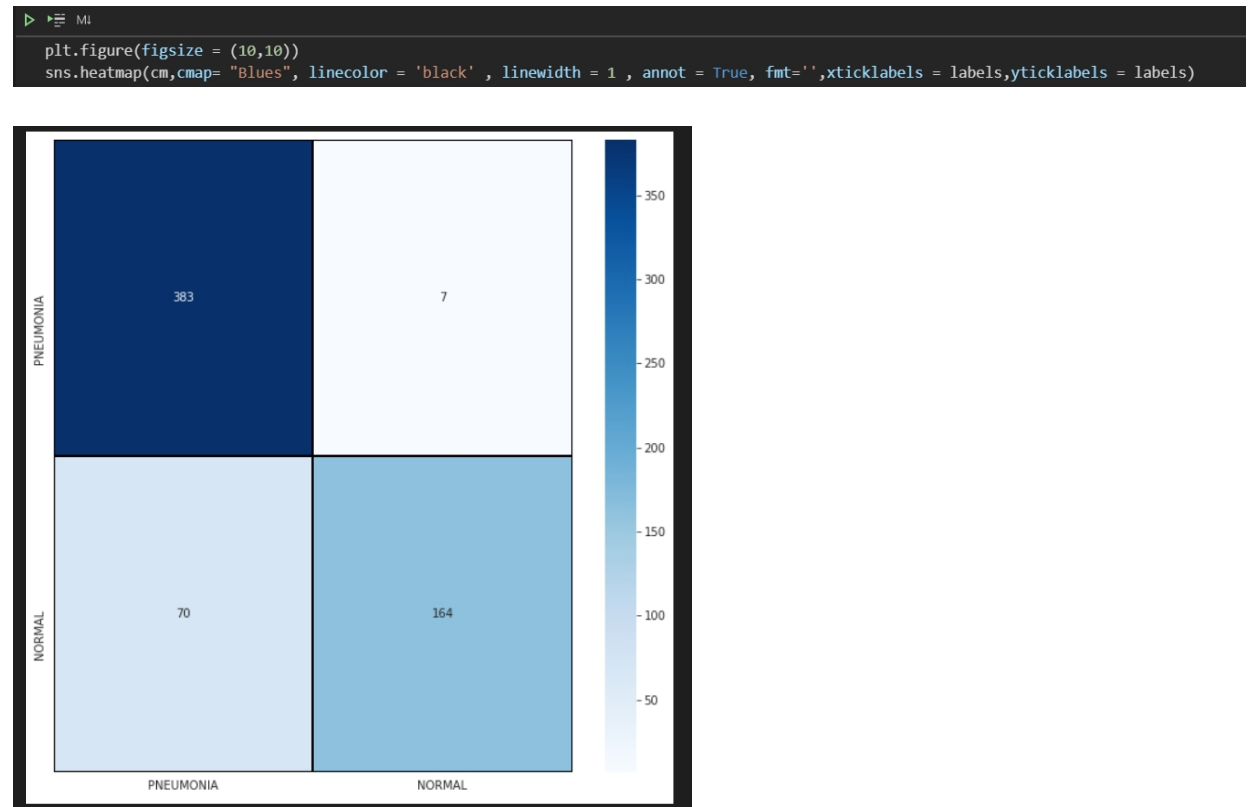
```
print(classification_report(y_test, predictions, target_names = ['Pneumonia (Class 0)', 'Normal (Class 1)']))
```

	precision	recall	f1-score	support
Pneumonia (Class 0)	0.85	0.98	0.91	390
Normal (Class 1)	0.96	0.70	0.81	234
accuracy			0.88	624
macro avg	0.90	0.84	0.86	624
weighted avg	0.89	0.88	0.87	624

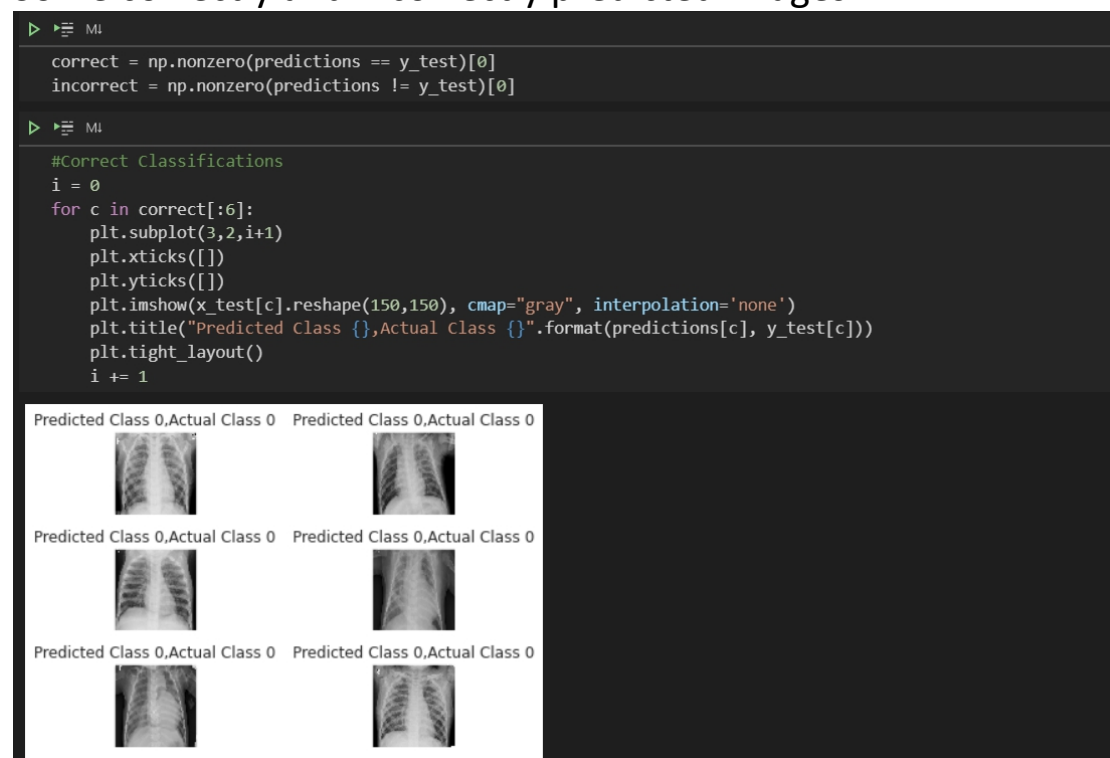
```
cm = confusion_matrix(y_test, predictions)
cm
```

```
array([[383,  7],
       [ 70, 164]])
```

Heat map :



Some correctly and incorrectly predicted images :





CONCLUSION :

Thus we have created a CNN model which can classify CXR images as normal or pneumonia with high accuracy.

The world of medical imaging is ripe for a revolution in terms of deploying CNN based technologies. There is no need for a doctor or a health care provider to ponder these images to gauge things. The task of reading these is incredibly menial and repititive. Those are two things that AI technologies are great at.

REFERENCES :

Dataset : <https://www.kaggle.com/paultimothymooney/chest-xray-pneumonia>

Conv Nets : <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

Python Implementation :
<https://www.datacamp.com/community/tutorials/convolutional-neural-networks-python>